



# Liabilities of AI-Driven Software Development and Guardrails in the Standard Model of Code

## 1. Documented AI Failure Modes (Anti-Patterns) – Validation & Refinement

AI coding agents exhibit several **failure modes** that differ from typical human mistakes. Below we revisit the **8 known anti-patterns** (e.g. *Context Myopia*, *Refactoring Abandonment*, etc.), validating or refining each with technical evidence and case studies:

- **Context Myopia:** AI agents often lack a holistic project view due to limited context windows, leading to code that doesn't integrate well beyond the snippet at hand <sup>1</sup> <sup>2</sup>. For example, Tenable researchers note that an LLM may suggest changes oblivious to other files, because “*usually the entire codebase will be larger than the LLM's context window*”, so it might miss relevant parts of the codebase <sup>3</sup>. Kodus similarly emphasizes that an AI “*doesn't have [the] contextual understanding*” of your specific system – it generates code from training patterns without grasping architectural requirements <sup>4</sup>. This myopia can cause **architecture misalignment** (code that violates layering or module boundaries) and **redundant implementations** of existing functionality. In practice, context myopia often manifests as duplicated helper functions or reimplementing logic already present elsewhere, simply because the AI didn't “see” it. A community analysis found that AI coding tools “*treat each request as a separate task*”, often failing to reuse context from earlier code, resulting in **redundant code** scattered across the codebase <sup>5</sup>. Over time, this increases maintenance burden and inconsistency.
- **Refactoring Abandonment:** AI agents may initiate refactors or broad changes but fail to propagate them throughout the codebase, leaving partial, broken transitions. A security blog on “vibe coding” warns that an AI often “*will suggest a refactor, but miss several places in the codebase that needed to be changed*.” <sup>6</sup> This leads to **inconsistent state** – e.g. a function renamed in some files but not others, or updated logic in one module but stale uses elsewhere. Developer anecdotes confirm this: Nolan Lawson notes coding agents love making “*subtle 'fixes' when refactoring code that actually break the original intent*”, such as inserting an extra null-check that wasn't needed <sup>7</sup>. These incomplete refactors create **semiotic mismatches** (names, comments, or tests that no longer align with code behavior) and can introduce bugs. The Standard Model's dimensions can help detect such rifts – for instance, a method's **D1\_WHAT** identity (name/structure) may diverge from its **D3\_ROLE** (intended purpose) after an abandoned refactor, signaling a “*Semiotic Misalignment*” <sup>8</sup>. Validation comes from practice: experienced engineers report using secondary AI “reviewers” to find these inconsistencies. For example, running a diff-checking agent on an AI-generated refactor often uncovers functional changes that should have been purely cosmetic <sup>9</sup>. This failure mode underscores why **thorough testing and diff-based validation** are critical whenever an AI performs a large-scale change.

- **Hallucinated Dependencies:** Generative AI may invent or suggest libraries, APIs, or modules that **do not actually exist**, a known phenomenon called *hallucination*. This is not just a harmless quirk – it creates a serious supply-chain liability. If a developer blindly tries to import the suggested package, they may end up pulling in a **malicious lookalike**. Endor Labs highlights this “*hallucinated dependencies*” risk: an AI might recommend a non-existent package, giving attackers an opportunity to publish a package under that name (a *slopsquatting* attack) <sup>10</sup> <sup>11</sup>. In one study, hallucinated imports were shown to potentially grant attackers full access to systems if the fake package is installed <sup>10</sup>. Kodus also notes that AI might reference “*non-existent libraries, methods, or classes*” because its training data had fragmented info <sup>12</sup>. To validate, researchers have measured a sharp rise in erroneous or vulnerable dependency usage linked to AI suggestions <sup>13</sup>. This failure mode is uniquely “AI-native” – human developers rarely type **import** statements for libraries that don’t exist, but AI will confidently do so. Mitigations include automated checks for unknown imports and **allow-lists** of approved packages (see §4).
- **Security Blind Spots:** AI-generated code frequently omits essential security steps or introduces subtle vulnerabilities. Studies show that developers using AI assistants are **more likely to inject security bugs** than those coding manually <sup>14</sup>. Common issues include missing input validation, weaker cryptography choices, or disabled error handling. In fact, over 40% of AI-generated code solutions were found to contain security flaws in one academic survey <sup>15</sup>. Many of these flaws align with known CWE Top 25 vulnerabilities, but appear in “AI-flavored” ways <sup>16</sup>. For instance, an AI might produce a perfectly functional feature that inadvertently *drops an authentication check* or uses an outdated, insecure API. Endor Labs researchers call this “*architectural drift*” – the AI’s code subtly breaks a security invariant without any obvious bug in syntax or logic <sup>17</sup>. An example would be replacing a secure default (say, parameterized SQL) with a riskier alternative (string concatenation) because the prompt didn’t specify security. Such changes often evade traditional static analysis because the code “*looks correct*” but violates an implicit assumption <sup>17</sup>. This anti-pattern highlights the need for **security-specific testing and linting** of AI contributions. It also suggests extending the Standard Model’s dimensions: e.g., an **Effect (D6)** analysis could flag if a function that previously had no external side effects suddenly writes to a new sink (potential data leak), or if an **Input/Output Boundary (D4)** is expanded in a way that bypasses validation layers.
- **Architectural & Layer Drift:** AI agents lack the higher-level understanding of a system’s architecture and design principles. They often pursue the prompt’s immediate goal in isolation, which can result in **violations of layering, modularity, or design patterns**. For instance, an AI might place business logic in a UI component or directly call a lower-layer service from a high-level module, skipping the intended abstraction layer. These *layer skips* and *role confusions* are akin to breaking a software “law of physics.” The Collider’s 8D classification can detect some of these: e.g., a function classified as **Controller** (an Application layer role) calling a **Repository** (Infrastructure layer) directly would violate the intended D2\_LAYER hierarchy. In fact, Collider’s built-in *antimatter laws* already define rules like *Layer Skip Violation (AM001)* to catch direct dependencies from e.g. a Controller to a database without a Service in between <sup>18</sup>. Evidence of this failure mode is found in AI’s tendency to produce “*short-term patchwork code*” without regard to overall structure. As AlterSquare reports, “*AI tools are great at solving specific problems but struggle with creating a well-organized system architecture*,” often treating each feature as standalone and accelerating technical debt <sup>19</sup>. Over time, this manifests as an incoherent topology: what might have been a clean layered *mesh* architecture degrades into a tangled *ball-of-mud* graph (e.g., cyclic calls between modules that should be separated). **Architectural drift** also includes AI inadvertently removing or bypassing

critical architectural elements – Endor Labs gives an example of an AI swapping out a cryptography module or “*removing access control protections*” during a simplification, thus undermining the system’s design goals <sup>20</sup>. Continuous architecture conformance checks are needed to catch these sneaky regressions (see §3 and §4).

- **Duplication and Redundancy:** A very common AI anti-pattern is **copy-paste coding** – the agent may duplicate logic in multiple places rather than refactoring or reusing existing code. Nolan Lawson wryly observes, “*I don’t know why agents love duplicating code so much, but they do. It’s like they’ve never heard of the DRY principle.*” <sup>21</sup>. This happens because the AI, lacking a global view, solves the same sub-problem anew each time it’s prompted, instead of abstracting it. Empirical analysis confirms this: AI-generated projects tend to have *more repeated code and boilerplate*. One startup report notes, “*AI tools frequently generate redundant code...leading to unnecessary duplication,*” such as multiple similar implementations of a feature where a human would write one reusable module <sup>5</sup>. Additionally, AI code can be *verbosely repetitive*: a function that a human might implement in 50 lines could balloon to 200 lines under AI, filled with redundant checks, overly defensive conditionals, or overly explicit conversions <sup>22</sup> <sup>23</sup>. These “*silly mistakes*” not only bloat the codebase, but also introduce divergence: if a bug is found in the duplicated logic, it has to be fixed in many places, increasing maintenance load <sup>24</sup>. In Standard Model terms, this anti-pattern erodes the **Composition (R4)** and **Relationship (R5)** integrity of the codebase – instead of clean connections between distinct components, you get parallel snippets doing the same thing. High duplication also confuses the **Ontology (R2)** lens, since multiple nodes represent what should be one concept. The remedy is to detect clones via static analysis (see §3) and enforce refactoring; indeed, projects using AI have started to set thresholds for duplicate code and reject AI contributions that don’t meet those standards <sup>25</sup>.
- **Semantic Drift & Documentation Mismatch:** Because AI doesn’t truly *understand* code semantics or the intent behind them, it can produce changes that drift away from the original meaning or fail to update documentation and tests accordingly. One facet is **Comment and Docstring erasure or obfuscation** – AI refactors sometimes strip out comments as “noise” or rewrite them incorrectly. Nolan Lawson gave an example where an AI turned a clarifying note like “this is a dumb hack we plan to remove in version X” into a polished but misleading comment, losing the original warning <sup>26</sup>. This is symptomatic of the AI’s lack of world-2 mental context (the rationale). Another facet is **purpose misalignment**: the AI might implement what it *thinks* the prompt asks, but in doing so, it can change subtle behaviors. For example, it might “fix” a bug by adding a null-check or altering a conditional, not realizing that breaks a downstream assumption <sup>7</sup>. Kodus refers to this as AI sometimes adding “*unnecessary logic... that wasn’t requested*”, which can impact performance or behavior in unanticipated ways <sup>27</sup>. Essentially, the code’s **meaning** drifts: function names, docs, or tests say one thing, but the new code does something else. The Standard Model’s **Role (D3) vs What (D1)** distinction is useful here – it can catch when a function’s name/category implies one role but its implementation now reflects another (e.g. a method still called `calculate_total` that after an AI edit also logs to a file, crossing from pure computation to I/O side-effect role). In fact, *semiotic misalignment* is explicitly flagged by comparing D1 and D3 in Collider <sup>8</sup>. This failure mode shows why adding an **Intent dimension (D9)** could be beneficial <sup>28</sup> <sup>29</sup> – modeling the intended behavior (from documentation or user stories) versus actual behavior can reveal when AI changes deviate from the spec.

- **Testing Gaps (Happy-Path Focus):** When AI is the primary developer, testing can be shallow or misdirected. AI-generated tests often give a *false sense of security* – they might achieve high coverage but miss edge cases and failure modes humans would catch <sup>30</sup>. Typically, AI-written tests focus on the “sunny day” scenarios that were described, ignoring error handling, boundary conditions, or integration aspects <sup>31</sup>. For example, an AI might verify a payment function works for a valid transaction but fail to consider an expired user session or network failure mid-transaction <sup>32</sup>. These blind spots mean critical bugs slip through until production. Moreover, if the AI is tasked with both writing code and tests, there’s a risk of **bias** – the tests may simply echo the code’s assumptions (or even be generated *from* the code), rather than truly validating it. This anti-pattern is evidenced by reports of AI-generated test suites that boast high coverage but still let serious bugs through <sup>33</sup>. It underscores that human insight is needed to design adversarial and exploratory tests. Within the Standard Model, this relates to the **Lifecycle (D7)** dimension – many AI failures occur in phases not explicitly specified (e.g., resource cleanup not implemented, error state not handled, etc.). Ensuring that each code unit properly addresses the full create→use→destroy lifecycle (D7) and using CI signals (like test failures or coverage drops) to flag anomalies is crucial (see §3 on detection).

**Summary:** These eight failure modes have been observed and corroborated by industry studies and community experience. They illustrate that AI agents, while powerful, have *natural liabilities* like narrow focus, lack of intent understanding, and no instinct for design consistency. The Standard Model’s multifaceted lenses help in diagnosing these issues – e.g., catching role mismatches, boundary violations, or dead code. In the next section, we introduce additional failure modes emerging uniquely from AI-as-developer scenarios, and then we detail how to detect and mitigate all these issues in practice.

## 2. Additional Failure Modes Unique to AI-First Development

Beyond the patterns above, when AI agents serve as primary developers, *new liabilities emerge* that are less common with human-only teams. Here are five (out of many) such failure modes, each grounded in recent observations:

- **Over-reliance & Skill Atrophy:** In an AI-centric workflow, human developers may become less vigilant, leading to a “*lazy mentality*” where they trust AI outputs without critical review <sup>34</sup>. Over time, this erodes deep coding skills (*skill erosion*), as devs no longer practice debugging, optimization, or architectural thinking <sup>34</sup>. The liability here is twofold: the team’s ability to catch AI mistakes diminishes, and the codebase accumulates design flaws that go unaddressed. This is a socio-technical anti-pattern unique to heavy AI use – the development team’s expertise “rusts” while the AI writes code. Organizations have noted this effect in surveys, warning that *continuous reliance on AI prevents developers from learning essential debugging and problem-solving skills* <sup>34</sup>. In the long run, this can leave a project extremely dependent on AI (and its limitations), with few engineers capable of understanding or evolving the system’s innards.
- **Non-Deterministic Changes (Heisenbug Code):** AI-generated code can introduce non-deterministic behaviors that are hard to trace. For example, an agent might generate a concurrent algorithm without properly handling race conditions, leading to flickering bugs that weren’t apparent in testing. Or, as some users reported with AI-assisted refactoring tools, the agent might apply changes inconsistently across runs – code that “works” in one generation may subtly change in a subsequent run. This unpredictability is related to the stochastic nature of AI generation. It’s a liability because it breaks the expectation of reproducibility in builds. While human developers can be inconsistent too,

AI can make this worse by, say, altering a function's behavior due to a slightly different prompt context. An anecdotal case involved an AI tool that, due to an internal state issue, would sometimes rename variables or alter logic differently on different days, causing confusion in code reviews (a true Heisenbug scenario). Mitigating this requires strict version control of AI prompts and outputs, and deterministically "freezing" AI-generated code once validated.

- **Excessive Dependency Introduction:** AI agents often solve problems by pulling in libraries or APIs (because the training data suggests them), where a human might write a small custom solution. This leads to **dependency bloat** – an AI might add numerous external packages, increasing attack surface and maintenance overhead. As one report noted, "*AI-generated code often includes unnecessary dependencies and overly verbose implementations*", such as using a heavy external API for a simple task <sup>23</sup>. In some cases, AI tools default to premium third-party services for convenience, incurring unexpected costs and complex vendor lock-in <sup>35</sup> <sup>36</sup>. For instance, an AI might integrate a cloud PDF parsing service where a local library would do, because it "saw" that pattern in training data. Over time, these extra dependencies become liabilities: they can introduce licensing issues, security vulnerabilities, or performance overhead. This is uniquely amplified in AI-driven development – a human is more likely to weigh the need for a new dependency, whereas an AI will merrily import anything that fits the prompt. Detection and mitigation involve scrutinizing new imports in PRs (see §3) and enforcing a "*contracts-first*" approach where the solution is defined before libraries are chosen (see §4).
- **Rapid Decay of Intent Traceability:** In an AI-generated codebase, understanding *why* a piece of code exists (the intent) can be particularly hard. AI code often lacks clear rationale or may not align with any design document, since it was produced from ad-hoc prompts. Humans usually leave behind clues (commit messages, design discussions, meaningful variable names) that indicate intent. AI, however, might produce code that is syntactically fine but cryptic in purpose. As a result, when requirements change or bugs appear, maintainers struggle to discern the original "thought process." This liability is essentially an extreme version of missing documentation. It's uniquely problematic in AI development because even the original author (prompt engineer) might not fully understand the code's logic (they accepted it because it worked). We see this in practice: teams have reported needing to hire senior engineers to "*untangle and rewrite code that junior developers can't manage*" due to lack of clarity in AI-generated logic <sup>37</sup>. This adds cost and delays. One way to address this is by enforcing that AI not only generates code but also a concise explanation or invariant for each major block (somewhat like a built-in documentation requirement – see mitigations in §4 and Manager Playbook in §7).
- **Cascade of Small Errors:** AI's mistakes can have a **compounding effect**. For example, an AI might introduce a tiny error (like off-by-one in an index) that a human would catch and fix immediately. But if left unchecked, the next AI prompt might adapt to that erroneous state, building more code atop it, and so on. We call this *error cascade*. In human teams, a bug would typically be caught in code review or testing before the next feature is built on top, breaking the chain. But an unsupervised AI agent could rapidly iterate, layering bugs upon bugs (especially if it's evaluating its own output). This results in a tangle that is far harder to debug because the origin is buried under subsequent changes. A case study in a continuous AI coding experiment showed that once a logic mistake slipped past initial tests, later AI revisions kept adjusting other parts of the program to accommodate it rather than eliminating it, leading to a convoluted design that ultimately had to be

scrapped. Such liabilities suggest that **feedback loops** and robust intermediate validation must be in place (so the AI corrects errors promptly instead of propagating them).

- **Model Misalignment with Evolving Codebase:** An AI agent's behavior is tied to its training and fine-tuning. As a project evolves (new frameworks, updated requirements), an AI might be *stuck in an outdated paradigm*. For instance, if an AI was fine-tuned on a codebase at the project's start, months later its coding style might conflict with the current architecture (because it's literally coding "from the past"). This can lead to friction where AI contributions consistently require rework to fit the new patterns – effectively an AI-induced legacy code problem. It's a liability because it slows down development (the AI drags the project backward). Humans learn and adapt, but a deployed AI model might need explicit updates. Until then, one can observe symptoms like the AI reintroducing deprecated functions or undoing agreed-upon refactors. This failure mode points to the need for continuous model updating or constraint of AI to narrow tasks if it cannot stay in sync with design changes.

The above additional failure modes – from human skill erosion to error cascades – underline that AI-first development isn't just "the same problems faster," but introduces qualitatively new challenges. Next, we discuss how we can *detect* these anti-patterns and anomalies in a codebase by expanding our analysis signals beyond simple timestamp forensics.

### 3. Detection Signals Beyond Timestamps

Historically, our team has used commit timestamps and temporal analysis (e.g. bursts of edits at odd hours) as clues for AI involvement or process problems. Here, we propose **10 concrete detection signals** that leverage static analysis, dependency graphs, CI results, runtime metrics, and topology analytics to catch the failure modes described. We also map each signal to the relevant Standard Model dimension(s) or lens:

1. **Duplicate Code Fragments (Static Analysis):** Measure the percentage of code that is highly similar (e.g. using a clone detection tool on each PR). A spike in duplicated blocks is a red flag for *Context Myopia* or *Duplication anti-pattern*. This correlates with **R4 Composition** and **R5 Relationships** – excessive clones mean poor composition and missing reuse relationships. A practical metric: *Clone Similarity Index* (e.g. any function with >30% of its code matching another function) triggers a warning. High duplication often indicates the AI didn't recall or generalize prior code [5](#).
2. **Unreferenced Code (Dead Code Introduction):** Track newly added functions or classes that are never invoked (no incoming "calls" edges in the dependency graph). If an agent introduces a function and nothing calls it, it might be an artifact of a hallucinated or abandoned plan. This signal ties to **D7 Lifecycle** (the code is created but never used) and **R5 Relationships** (no relationship edges to the node). Forensic query: find nodes with `calls_in = 0` that appeared in the last commit – these could be orphan "ghost" functions indicating incomplete integration.
3. **Cyclomatic Complexity & Size Spike:** Compute cyclomatic complexity and function length for new/changed code. AI-generated code that is overly complex or lengthy for its purpose may indicate *Silly Mistakes* or *Over-engineering*. For example, if a simple controller method jumps from a complexity of 4 to 20 after an AI edit, that's suspicious. Map this to **D1 WHAT** (if the atom type is "Function" and it violates complexity norms for its type) and to the **Quality lens** (if defined). We can set thresholds

(e.g. no function > X lines or complexity Y without manual review). Complexity metrics from information theory can be applied here (e.g. Halstead volume, or Shannon entropy of token distribution <sup>38</sup>) – an anomalously high entropy or volume for code that nominally does a simple task signals a likely AI artifact.

4. **Layer Violation Edges:** Analyze the dependency graph for any new edges that violate layering constraints (e.g., Infrastructure layer calling Application layer, or skipping layers). Collider can directly flag these via antimatter rules (AM002 Reverse Layer Dependency, etc.) <sup>39</sup>. A new instance of such an edge appearing post-commit is a strong sign of *Architectural Drift*. This is mapped to **D2 LAYER** and **D4 BOUNDARY** dimensions. A detection query might be: *for any new import or call edge, check if  $L(\text{source}) < L(\text{target})$*  – if yes, we have a *reverse dependency* which is usually forbidden <sup>40</sup>. Similarly, check for layer skips (calls that jump more than one layer distance).
5. **Unapproved Dependency Introduction:** Integrate a check in CI for any new external library/package added in a commit. Cross-reference new imports with a whitelist or with public package indexes to see if they exist. This catches *Hallucinated Dependencies* (nonexistent ones) and flags potentially risky ones. It aligns with **D4 BOUNDARY** (crossing into an external boundary) and **R5 Relationships** (new external relationship). For example, if a PR adds `import foobar` and `foobar` is not in the project's dependency file or known package list, raise an alert. This signal directly combats the slopsquatting scenario <sup>10</sup>. Even for real packages, if an AI adds a dependency that duplicates existing functionality, it could be flagged for review (mapping to **D3 ROLE**: is the AI introducing a new utility that overlaps with an existing one?).
6. **Static Security Scan Delta:** Run a static application security test (SAST) on the codebase before and after AI changes. New **security findings** (e.g. an XSS vulnerability or use of a weak cryptographic function) attributed to the latest diff indicates the AI may have injected a flaw. This covers things like missing input validation, use of `eval`, hardcoded credentials, etc. Map this to **D6 EFFECT** (side effects and external interactions) and **D4 BOUNDARY** (since many security issues occur at input/output boundaries). For example, if an AI adds a file upload feature and the SAST now flags “unsanitized file name used in file path” – that’s a concrete signal. We treat any increase in high-severity static analysis warnings as a must-review for the AI changes <sup>41</sup>.
7. **Test Coverage and Quality Signals:** Monitor test coverage % and test outcome patterns. If an AI-generated PR significantly reduces overall coverage (e.g. lots of new code without tests) or if the pattern of test failures is unusual (many edge-case failures, or tests passing because they are too trivial), this is a signal. Specifically, **lack of edge-case tests** can be detected by analyzing test code: e.g., are all tests using only well-formed inputs? (AI tests often ignore “messy” inputs <sup>31</sup>). This relates to **D7 LIFECYCLE** (ensure tests cover create/use/destroy phases) and **R8 Epistemology** (confidence). A practical approach: require that any new function above a complexity threshold comes with corresponding new tests (enforced via CI). If not, label it potential AI oversight. Also, fuzz testing results can be a signal – if a component written by AI suddenly fails fuzz tests or property-based tests that similar human-written components pass, it indicates blind spots.
8. **Runtime Telemetry Regression:** In a continuous delivery setup, compare key runtime metrics before vs. after a deployment that includes AI-authored code. Signals include a jump in error rates (e.g. exception frequency), memory leaks, or latency degradation in specific modules. For example, if after an AI-generated data processing module is released, the memory consumption per request

doubles or you see intermittent timeout errors, this quantifies a *Performance/Resource issue* possibly introduced by AI (maybe due to inefficient algorithms or missing cleanup). Map this to **D6 EFFECT** (the effect on execution) and possibly a new lens like *Operational lens*. Automated anomaly detection on metrics (using thresholds or ML) can catch these. Each anomaly should be traced back to the code changes that could cause it. Many AI mistakes (like  $O(n^3)$  loops or not closing file handles) will surface here if not caught in review.

9. **Topology Changes in Collider Visualization:** Use Collider's graph outputs to spot unusual topology patterns after AI contributions. Signals might be:
10. **Star-shaped clusters** – e.g., one new “God function” that everything calls (could indicate the AI created a mega-utility or *God Class* anti-pattern). This reflects in **Topology metrics** like *knot\_score* or *centrality measures*. A high increase in centrality of a node is mapped to the idea of a *Pathological Holon (God Class)*.
11. **Increase in isolated nodes or micro-islands** – code that is not well integrated (possibly dead code or duplicate logic not wired in). This can be seen in **reachability\_percent** dropping or number of weakly connected components increasing <sup>42</sup>.
12. **Layer entropy** – measure if a file’s content now spans more dimensions than before (e.g., functions of many different roles all in one file – sign of muddled design).

These signals tie to **R5 Relationships** and **R7 Semantics** lenses. Essentially, if the qualitative shape of the dependency graph changes for the worse (more spaghetti edges, more hubs, or more disconnected pieces), it’s a red flag the AI introduced structural issues. Collider’s *visual topology* analysis (star, mesh, etc.) <sup>43</sup> can be automated: e.g., if a codebase was a nicely layered mesh and after a PR it trends toward a chaotic star/hub, mark that PR.

1. **Dimension Drift Tracking:** Track the Standard Model dimension classifications of code over time. If a particular function’s classification in certain dimensions changes unexpectedly from one commit to the next, that might signal an AI-induced misclassification. For instance, suppose a function was consistently categorized as **Internal** boundary (D4) but after an AI edit it’s detected as **I/O** boundary – that suggests the function now does I/O it didn’t before (maybe the AI added a file write or network call). Similarly, a method might flip from **Stateless** to **Stateful** (D5) or from a pure **Utility** Role to a mixed Role. These drifts can be automatically detected by running Collider classification on each version. If a dimension change doesn’t align with an intentional design change, it could be an AI anti-pattern surfacing. In practice, one could maintain a log of each file’s dimensional signature (D1–D8) and alert if any dimension (except D8 trust/confidence) changes beyond a threshold without corresponding justification (e.g., a pure function becoming impure should prompt a review). This leverages **all dimensions**, especially D2–D6, and maps to the concept of maintaining a consistent taxonomy. A sudden **D3\_ROLE ambiguity** (e.g., a new function that the model can’t confidently classify into one role) is also a sign – it might be doing too much (violating single responsibility), a hallmark of AI’s “kitchen-sink” methods.

Each of these signals can be integrated into CI/CD pipelines or analysis dashboards. By going beyond timestamps (which only told us *when* something suspicious happened) to examine *what* the code is doing, we get a multidimensional view of AI-induced issues. Many signals correspond to known metrics (complexity, coupling, etc.) and can be mapped back into the Standard Model’s schema for automated “health checks.” In fact, the **Collider’s semantic space** is well-suited to define these as queries or additional lenses (e.g., a “Violation lens” that projects any anti-pattern signals onto the nodes).

In the next section, we turn to mitigation strategies – how can we proactively prevent or reduce these failure modes? We will expand on existing strategies (like Linter Laws, Contracts-First, Socratic Supervisor) and propose new ones, with concrete steps for implementation.

## 4. Mitigation Strategies and Best Practices

To guard against the above liabilities, we need robust **mitigation strategies**. The key is not just high-level policies, but *operationalized practices* – specific instructions for AI agents, enforced rules in CI, and process changes that are testable. Below we detail the three provided strategies and add at least three more, each with concrete implementation guidance:

- **Linter Laws (AI Guardrails via Static Analysis):** This strategy treats coding standards and architectural rules as inviolable “laws” that both AI and humans must obey. In practice, this means encoding style guides, design patterns, and the Standard Model’s invariants into linter or static analysis rules, and making the AI aware of them. For example, we can extend ESLint or Pylint with custom rules reflecting our architecture (e.g., “Controllers must not directly import Repositories” or “No usage of `eval()` allowed”) <sup>44</sup>. These become **pre-merge checks** in CI – any violation blocks the PR. To operationalize for the AI agent: include a summary of these Linter Laws in the agent’s system prompt or `CLAUDE.md` (the project context file) <sup>45</sup>. E.g., “*Your code must pass all style and layer checks: see `antimatter_laws.yaml` for forbidden patterns.*” Additionally, after the AI writes code, run an automated linter pass. If issues are found, either auto-correct them or prompt the AI to fix them before human review. This was inspired by the idea of a “Constitution” for the AI – effectively, encode the project’s do’s and don’ts as laws. The Collider tool’s schema can integrate these as well; for instance, listing current *violations* of antimatter laws in the Brain Download report so developers and the AI can see them. **Example:** If Linter Law says “No function over 100 lines,” and the AI produces one, the CI rejects it with a clear error. The AI, seeing this feedback (if integrated), should then split the function or simplify it. By enforcing such laws consistently, many AI slip-ups (style issues, obvious anti-patterns) get nipped early. Over time, as the AI is tuned or few-shot trained on our `antimatter_laws.yaml`, it will likely produce fewer violations from the outset.
- **Contracts-First Development:** This strategy mandates that development starts from **specifications and interfaces** (contracts) before implementation. In an AI context, that means using the AI to first generate things like function signatures, type contracts, API schemas, or even tests – *without writing the core logic yet*. For example, one might prompt the AI: “*Design the interface and docstring for a function that does X, following our conventions, and propose some usage examples/tests.*” Only after the contract is reviewed (by a human or a supervising agent) do we let the AI fill in the implementation. Contracts-first mitigates misalignment by anchoring the AI to an intended shape. It also helps address Context Myopia: if we have a clear contract, the AI is less likely to drift outside its boundaries. **Operationalization:** Define templates or schemas for common contracts (e.g., function signature with pre/post-conditions, class interface, or an OpenAPI spec for a service endpoint). Incorporate a step in the development pipeline where the AI must output these templates. The *Socratic Supervisor* (next strategy) or a human reviews the contract for sanity – ensuring it doesn’t violate any design (like correct layering: e.g., if a function is meant for the Core layer, its contract shouldn’t involve UI concepts). In CI, you could enforce that for each new module, an interface spec (or at least a docstring with expected behavior and constraints) is included – essentially, fail the build if a new public method lacks a contract. This approach draws from design-by-contract and test-driven development: some teams have AI generate unit tests first as a form of contract, which then

guides the implementation. By focusing the AI on “*what*” and “*for whom*” before the “*how*,” we reduce off-track implementations and ensure important considerations (like performance budgets or security requirements) are declared up front. For instance, a contract for a data processing function might specify “must handle up to 10k records and maintain O(n) complexity” – the AI will then be aware of that constraint while coding.

- **Socratic Supervisor (Interactive Oversight):** This strategy introduces a supervisory agent or role that continuously interrogates and audits the AI’s decisions using a Socratic method. Instead of the AI operating in a vacuum, the Socratic Supervisor (which could be a human facilitator or another AI configured to critique) asks pointed questions at each step: “*Why did you choose this approach? Could there be edge cases? Is this consistent with module X’s design?*” The goal is to surface hidden issues and force justification. Nolan Lawson’s technique is a great real-world example: after letting the AI propose a solution, he’d start a new session and prompt, “*Review this diff as if you are a senior engineer – did the changes introduce any functional bugs? Are they truly a pure refactor?*” <sup>9</sup>. The AI acting as a reviewer often finds the mistakes of its “coder” counterpart. To operationalize Socratic supervision, one can set up a two-agent system: **Coder AI** and **Critic AI**. The Critic has a prompt like “You are a code reviewer with 20 years experience. You will receive code diffs and point out any potential issues, ambiguities, or departures from best practices.” Every time the Coder produces something non-trivial, the Critic agent is invoked to analyze it (this can be done asynchronously in CI or during the agent conversation). Importantly, the Critic should be empowered to halt the process if certain red lines are crossed. For instance, if the Coder writes a DB query without using parameterization, the Critic should flag a SQL injection risk immediately. In practice, tools implementing this pattern have shown significantly higher success rates – the AI catches its own errors when forced to explain or defend them <sup>9</sup>. Within the Standard Model context, we can think of the Socratic Supervisor as providing an **Epistemology lens (R8)** feedback loop – it challenges the certainty of the AI’s answers and asks, “How do we know this is correct?”. One could even integrate this into Collider by having a mode where for each Brain Download insight or each recommended change, a set of heuristic questions (derived from our anti-pattern list) are asked and answered. The Socratic method ensures continuous validation of assumptions, mitigating issues like hidden security flaws, misinterpretation of requirements, and logical errors that an uncritical agent would miss.
- **Test-Driven AI Development (Continuous Testing Hooks):** This mitigation extends the Contracts-First idea into execution: require that AI-written code is accompanied by automated tests, and leverage CI to enforce that tests (both new and existing) pass at each step. Concretely, before an AI commit is accepted, the system runs all tests in a sandbox. If any regression is detected, that feedback is immediately given to the AI to correct the code. This creates a tight feedback loop, akin to having a safety net. To operationalize: integrate a testing harness in the agent’s environment. Some AI coding agents can run `pytest` or similar on command; you can script the agent to run tests after writing code (or even have it write a test for each new function). If a failure occurs, the agent should attempt to fix it (this is similar to how tools like Ghostwriter or Replit’s AI can iteratively debug). This practice addresses *Happy-path bias* by actively forcing consideration of failing scenarios. It also mitigates *error cascades* – since errors are caught as soon as they are introduced, the AI is less likely to build on faulty code. One can formalize this as a CI **gate**: e.g., “*All functions changed by the AI must have at least X% coverage and all tests (old and new) must pass.*” The Standard Model’s D7\_LIFECYCLE can be extended to include a “Tested” status (perhaps as part of a lens or an annotation on each particle: tested or not). With test-driven prompts, you shift the AI’s focus from “just make it work” to “make it correct.” Empirical evidence shows AI agents perform better when

given unit tests to satisfy – it turns coding into a puzzle where the tests define success, reducing ambiguity.

- **Multi-Agent Pair Programming:** Introduce diversity and specialization by pairing AI agents with distinct roles – e.g., one agent writes code, another focuses on architecture enforcement, another on performance review. Each agent has different prompt conditioning (one might have a prompt emphasizing “be concise and optimize for readability”, another “be strict about architecture and layer rules”). The idea is to simulate the checks-and-balances of a dev team within the AI itself. For instance, after the coding agent proposes a solution, a “Reviewer agent” can be prompted (similar to Socratic Supervisor but perhaps more focused on static aspects) and a “Tester agent” can try to break the code (generating adversarial inputs or property tests). They then discuss or provide feedback to a coordinator that decides the next action. This approach mitigates single-agent blind spots – an agent focused on speed might miss security, but the security-focused agent can catch it. It’s operationalized by orchestrating the agents via a simple script or workflow: code -> review -> test -> iterate. One could use a tool like an AI planning framework (e.g., LangChain or a custom harness) to manage this. The Standard Model’s *8D manifold* can even be split among agents: e.g., assign one agent responsibility for ensuring D2 and D3 (layer and role consistency), another for D4 and D6 (boundary and side-effect concerns), etc., effectively having each agent “own” a subset of dimensions to watch. To test this practice, one could compare outcomes of a single agent vs. multi-agent on complex tasks – we expect fewer anti-patterns slip through the multi-agent process. The overhead is higher (compute cost, complexity of coordination), but for critical code it might be justified.
- **Continuous Quality Monitors and Evolution Protocols:** This strategy is about **sustained vigilance** on the codebase health over time, not just per-PR. It involves setting up automated monitors for the detection signals (from §3) and defining an *evolution protocol* – when a metric goes out of range, there is a planned intervention. For example, if code duplication (per Collider’s analysis) exceeds, say, 5%, you schedule a refactoring sprint (perhaps AI-assisted) to DRY it up. Or if the coupling between certain modules grows beyond a threshold, you task the AI (under supervision) to reorganize or modularize. Essentially, instead of waiting for a human to notice the system rotting, the monitors trigger tasks proactively. This is mitigation in the sense of limiting long-term damage. Operationally, one could integrate Collider outputs into a dashboard that managers and leads review regularly. The **8D model** provides dimensions to focus these monitors – e.g., track D5\_STATE distribution (maybe too many global stateful singletons cropping up could be an AI anti-pattern – time to enforce some state handling rules) or D8\_TRUST (if classifier confidence is dropping for a lot of new nodes, maybe the new code is confusing or inconsistent). The *Collider’s antimatter\_laws.yaml* can be extended with more laws as we identify new anti-patterns, and the continuous analysis will flag violations. Managers can then enforce that “no critical violations are outstanding before a release.” One could even codify some of this into a **CI enforcement logic**: for instance, fail the build if a new commit raises the “God Class” violation or if it pushes the *dead\_code\_percent* beyond a set point <sup>46</sup> <sup>47</sup>. By tying these monitors to concrete actions (e.g., create a ticket for addressing each violation, auto-assign to the AI or a dev), we close the loop between detection and remediation – ensuring mitigations are not one-time, but ongoing.

Each of these strategies should be **concrete and testable**: we can simulate scenarios to verify they work. For instance, intentionally introduce a layer violation and see if Linter Laws catch it in CI; or have the AI attempt a risky refactor and observe if the Socratic Supervisor flags the potential issues (we could script an internal test where the Critic agent must catch a planted bug). Moreover, these practices are meant to be

integrated: e.g., Linter Laws feed into both the AI's prompting (so it knows the rules) and the continuous monitors (so any violation is auto-detected).

We have now a set of anti-patterns and mitigations. The final step is to weave these into our **8D classification model and Collider schema**, to ensure our tools and ontology can formally represent and check these issues.

## 5. Integration into the 8D Model and Collider Schema

To systematically manage AI failure modes, we should extend our existing **8D Standard Model and Collider tooling**. This involves mapping the new insights to dimensions, possibly introducing new dimensions or lenses, and updating Collider's schema (e.g., the `antimatter_laws.yaml` rules or `violations` fields) to codify these patterns. Below is a plan for integration:

### Mapping Failure Modes to 8D Dimensions

Each failure mode can be linked to one or more Standard Model dimensions (D1–D8), indicating where in the manifold the issue manifests:

- *Context Myopia*: Manifests in **D2 Layer** and **D5 State**. The AI's lack of context often causes it to violate architectural layering (using wrong layers) and to mishandle state (e.g., reinitialize something that should be single-instance). It also shows up in **R5 Relationships** – missing references to existing modules. **Lens mapping**: The *Hermeneutics lens* (R7) could be useful to detect context issues by interpreting code in context of file paths; context myopia often means the code's location and content mismatch (e.g., a utility function doing domain logic because the AI didn't see the domain module). We might consider a new lens for "Contextuality" that assesses how much of the project context was considered (hard to quantify, but perhaps by looking at the breadth of imports or references).
- *Refactoring Abandonment*: Appears in **D7 Lifecycle** (incomplete lifecycle of a change – something created but not properly updated or destroyed elsewhere) and **D1 What vs. D3 Role** (signatures/names that no longer match roles). Also ties to **D8 Trust** – after an abandoned refactor, our confidence in those portions should drop, which Collider could reflect by a lower trust score if it detects inconsistencies. **Lens**: *Epistemology* (R8) lens is apt here – highlighting areas of low certainty which might correspond to half-finished changes. We could enhance Collider to identify sudden drops in confidence or heuristic signals that a change was not uniformly applied (for example, multiple functions with similar names except one differs – could be a missed refactor spot).
- *Hallucinated Dependency*: Mostly a **D4 Boundary** issue (calls going out to external libraries that might not exist or should not be used). It also hits **D3 Role** if the AI pulled in a dependency that doesn't fit the intended role (e.g., using an external service for something meant to be an internal utility). **Lens**: *Relationships* (R5) lens can highlight external vs internal calls; we might refine Collider to mark nodes that correspond to external packages and flag if they're suspicious (maybe by checking if they appear in known package manifests). Perhaps a new *SupplyChain* lens could be introduced, focusing on dependency health (e.g., using metadata on package popularity or known vulnerabilities).

- *Security Flaws*: These cut across dimensions but often involve **D6 Effect** (side effects and I/O not handled safely) and **D4 Boundary** (improper input/output handling). If we had an **Intent dimension (D9)** to capture intended security level or invariants, misalignment there would catch it (for example, if a function is intended to enforce authentication but its code doesn't). **Lens**: Transformation (R6) lens (what does it do, input → output) can be enhanced to check for missing validation or dangerous transformations. Possibly incorporate a “taint” analysis within R6 to see if inputs flow to outputs without checks.
- *Architectural/Layer Drift*: Directly mapped to **D2 Layer** and secondarily **D3 Role**. E.g., if an AI moves some logic to a different layer or mixes roles, Collider should detect an element whose D2 or D3 classification is out-of-place relative to its file location or naming. We might also use **R4 Composition** lens: architecture drift often results in huge classes or god objects (which Composition lens would reveal as an outlier in number of parts). If needed, a *D9 Intent* could capture architecture intent (like marking certain modules as “should only contain Core layer atoms”) so drift could be measured as deviation from that intent.
- *Duplication & Redundancy*: Not a single dimension, but can be recorded in **R5 Relationships** (lack thereof, since duplicates don't call each other) and a potential **R9 lens** if we had one for *Quality*. Alternatively, we add a field or metric in Collider's output for duplication clusters. Since duplication is like an “antimatter” that violates the ideal model, we might treat it as a violation without needing a dimension (similar to how we treat God Class as an anti-pattern rather than a dimension value).
- *Semantic Drift (Doc/Name vs Behavior)*: This is exactly the gap that adding **D9 INTENT** would fill, as previously discussed <sup>28</sup>. We should integrate documentation and intent into the model. For instance, D9 could be an “Intent” dimension with values like “Documented vs Implicit vs Contradictory” <sup>48</sup>. A failure mode here is when code is *Contradictory* to its docs – Collider can flag that as a D9 discrepancy (with detection via NLP comparing docstring to code behavior <sup>29</sup>). Also, *semiotic misalignment* can be seen by comparing **D1 What** (e.g., a class is named “Manager”) with **D3 Role** (Collider infers it's doing something else) <sup>8</sup> – that is effectively a semantic drift signal.
- *Testing Gaps*: This would benefit from introducing a new lens or field capturing test coverage or validation status. Perhaps we add a **“Verification” lens** that for each function can note if there are corresponding tests or not, or even link to test cases that hit it (if we parse coverage reports). Not originally in 8 lenses, but could be an extension: an R9 Verification lens asking “How is this verified?” with answers like “Not covered / Unit tested / Property tested / Proven”. For now, we integrate test signals into Collider by enriching the node metadata with a flag if it's mentioned in test files or its behaviors are asserted. This directly addresses the AI's testing blind spots by making visible which parts of code are untested.
- *Skill Erosion (Team aspect)*: This one is external to the code itself, so not directly mappable to 8D. However, we could introduce a **Process dimension** or track involvement metrics (like how often code is AI-authored vs human-authored). Perhaps an additional metadata field in nodes or the meta section could be “authorship: human/AI” if we want to store that (though detecting that reliably is non-trivial). Another idea is a **Lens of Evolution**: how code came to be (via human or AI), which might be logged in an external ledger. For the Standard Model, we might keep this out-of-band but still crucial for managers to monitor (e.g., if 90% of code in a subsystem is AI-written, extra scrutiny might be applied).

In summary, most failure modes map to existing dimensions or highlight the need for **D9 Intent** (for semantic alignment) and possibly **D10 (Language/Platform)** if some issues are language-specific (already proposed <sup>49</sup>). The existing 8D system is largely sufficient, but adding **Intent** and maybe a **Quality/Test lens** would provide first-class support to detect these liabilities.

### Enhancements to Collider Schema (`antimatter_laws.yaml` and `violations`)

Collider's `antimatter_laws.yaml` currently houses rules for layer violations and god classes <sup>18</sup>. We should extend this taxonomy of "antimatter" (destructive patterns) to cover the new failure modes:

- **Duplicate Code Law:** Add a rule to detect significant copy-paste overlap. For example:

```
- id: "AM004"
  name: "Duplicate Logic"
  description: "Similar or identical code appears in multiple places."
  detection: "If two functions >30 lines share >80% similarity and are not just overrides."
  severity: "WARN"
  examples:
    - "Two validation functions in different classes doing the same checks with minor differences."
    - "Copied and pasted utility logic instead of reusing."
```

Collider (or a pre-processing step) can compute hashes or fingerprint of functions to find duplicates. Violations would be attached to each duplicate node.

- **Orphan Function Law:**

```
- id: "AM005"
  name: "Orphan Code"
  description: "Code implemented but never used (potentially dead or forgotten feature)."
  detection: "Function has no incoming call references in the project."
  severity: "INFO"
  examples:
    - "Helper method added by AI that nothing invokes."
```

This leverages the graph: any node with `in_degree(calls) = 0` (and isn't an entry point or test) triggers it. While dead code isn't always urgent, it is a sign of possible AI oversights.

- **Hallucinated Import Law:**

```
- id: "AM006"
  name: "Unknown Dependency"
```

```

description: "Introduction of an import that is not recognized or allowed."
detection: "New import not present in package metadata or allow-list."
severity: "ERROR"
examples:
  - "Importing a module that is not installed (likely hallucinated)."
  - "Using a banned library (e.g., deprecated or insecure dependency)."

```

To implement, Collider's analysis could consult a config of known dependencies and mark any foreign ones. At minimum, this can be part of CI gating.

- **Contract Violation Law:** If we formalize certain API contracts (say via function annotations or interface schemas), we could add laws to catch when code violates them. For instance, if a docstring says `#@pure` but the function does I/O, or if a function is meant not to throw exceptions but does. This might require more advanced static analysis or using the Intent dimension. We can leave this as a future hook, but mention it in the schema as a placeholder for "Specified Intent vs Implementation mismatch".
- **Test Coverage Law:** If we integrate test data, a rule could be: any new function above X complexity with zero tests is a violation. This again might be fed by an external coverage tool.

These antimatter laws serve as automated "gotcha" detectors in Collider's output. We should also add a `violations` field in each node's JSON output to list which laws it breaks (if any). For example:

```
{
  "id": "src/util/math.py::calculate_total",
  "name": "calculate_total",
  "dimensions": {...},
  "lenses": {...},
  "violations": ["AM004", "AM005"]
}
```

This would mean `calculate_total` has been flagged for Duplicate Logic (maybe it's nearly identical to `compute_sum` elsewhere) and Orphan Code (perhaps nothing calls it yet). The HTML report could highlight these in red, guiding reviewers to suspicious spots.

We may also consider if new **dimensions or lenses** are needed. As discussed, **D9 Intent** is a strong candidate to add explicitly <sup>28</sup> – it helps capture what the code was *supposed* to do (via documentation or naming) versus what it does. We might also introduce a **Lens for Process or Provenance** that isn't one of the original eight lenses: something like "Who/what created this code? Under what assumptions?" – this could store metadata like commit author, or a simple AI-vs-human flag if detectable. This lens could be immensely valuable in post-mortems (e.g., if an issue is found in production, knowing it came from an AI-generated chunk might shape the debugging approach).

Additionally, Collider's **visual topology** can be extended with an overlay for anti-patterns. For instance, highlight nodes with violations in a distinct color or shape (like a  $\triangle$  icon on nodes that have antimatter law hits). This makes the *Collider visual map* not just a neutral view but an actionable hotspot map.

Finally, we incorporate these into our **8D classification system documentation** and training. The anti-pattern taxonomy (which was dynamic per `LEARNING_LEDGER.md`) will grow with these new entries, and the model (if it's an AI model classifying) can be trained on recognizing them. For example, patterns like "function that only differs in one constant from another function" could become a learned pattern for duplication.

In summary, by mapping each failure mode to the Standard Model's facets and extending Collider's rules, we ensure that our theoretical understanding translates into concrete checks. The 8D model remains our "coordinate system" for code, and now these failure modes become recognizable constellations in that system – detectable, measurable, and ultimately fixable.

## 6. Example Metric Tables and Schema Snippets

*(As part of the integration, we provide a brief example to illustrate how some metrics and rules could be documented for our team's reference.)*

**Table 1: Proposed Detection Signals and Mapping**

Signal	Description	Standard Model Mapping	Example Threshold
Duplicate Code (% similarity)	Measures code clone redundancy in new commits.	R4 Composition, R5 Relationships	>80% similarity over 30 LOC.
New Dependency Count	New external imports added by a commit.	D4 Boundary, R5 Relationships	>0 (flag each new external import).
Cyclomatic Complexity	Max complexity of new/ changed functions.	D1 What (Function), Quality lens	>10 for any single function.
Layer Violation Edges	Cross-layer calls/imports that break architecture.	D2 Layer, D4 Boundary	Any instance triggers ERROR.
Security Warnings (SAST)	New high-severity static analysis alerts.	D6 Effect, D4 Boundary	Any instance triggers FAIL build.
Test Coverage Delta	Drop or insufficient tests for new code.	D7 Lifecycle, (proposed R9 Verify)	<70% cov on new functions.
Dead Code (Orphan)	New code with no usages (in-degree 0).	D7 Lifecycle, R5 Relationships	Any orphan function - > WARN.
Performance Regression	Change in runtime metrics post-deploy.	D6 Effect (exec), R6 Transformation	>10% slowdown or errors +X.

Signal	Description	Standard Model Mapping	Example Threshold
Role Drift	Change in inferred Role (D3) for an existing entity.	D3 Role, D1 What	Role change without schema update.
Intent Mismatch (Doc vs Code)	NLP diff between docstring intent and code behavior.	D9 Intent (proposed), R7 Semantics	Any contradiction -> WARN.

(Table above lists some signals, their mapping to the model, and example criteria for flagging.)

### Snippet: Antimatter Laws Extensions (partial)

```

antimatter_laws:
  - id: "AM001"
    name: "Layer Skip Violation"
    ... # (existing law for skipping layers) ⑯

  - id: "AM002"
    name: "Reverse Layer Dependency"
    ... # (existing law for reverse dependency) ⑯

  - id: "AM003"
    name: "God Class (Pathological Holon)"
    description: "Class with too many responsibilities, indicated by excessive size or mixed roles."
    detection: "If class > N LOC and implements > M distinct Roles (D3) or has > K methods."
    severity: "WARN"
    examples:
      - "A single class implementing UI, business logic, and data access."

  - id: "AM004"
    name: "Duplicate Logic"
    description: "Substantial code duplication detected across functions."
    detection: "If Jaccard similarity of AST between two functions > 0.8 over >30 LOC."
    severity: "WARN"
    examples:
      - "Two functions in different modules performing the same calculation with minor differences."

  - id: "AM005"
    name: "Orphan Code"
    description: "Function or class is never referenced (potential dead code)."
    detection: "No incoming calls or usage references in the project (excluding tests)."

```

```

severity: "INFO"
examples:
  - "A helper function added by AI that nothing calls."

- id: "AM006"
  name: "Hallucinated Dependency"
  description: "Import of a non-existent or disallowed external package."
  detection: "New import not found in known dependencies or in stdlib list."
  severity: "ERROR"
  examples:
    - "import foobarbaz (package not installed or typo in name)."
    - "from old_crypto import AES (using deprecated insecure lib)."

```

With these rules in place, Collider's report would, for instance, list "AM004 Duplicate Logic" next to functions that match, helping reviewers to quickly identify and consolidate duplicate code.

## 7. Manager Playbook for Enforcing AI Guardrails

Finally, for non-technical project leaders, here's a concise playbook to ensure these guardrails are in place and effective. These are management-level practices to *steer the team and process*:

- **Define "AI Do's and Don'ts" Upfront:** Establish a lightweight policy document (in plain language) that highlights forbidden practices (e.g., "Don't accept code that introduces new dependencies without approval", "AI-written code must follow our style guidelines"). Make sure the team and any AI tools are aware of this. As a manager, insist these rules are discussed in planning and retro meetings.
- **Implement Quality Gates in CI:** Work with the dev team to add automated checks (linters, tests, Collider analysis) to the pipeline. As a lead, you don't need to code these yourself, but ensure they exist and have **teeth** (e.g., the build fails if an AI anti-pattern is detected). Ask for a dashboard of key metrics (duplication rate, coverage, etc.) and review it regularly. This makes quality **visible** and non-negotiable.
- **Enforce Pair Review (Human or AI):** Never allow AI code to go to production unreviewed. If you lack enough senior engineers, consider using an AI code review tool (like Kodus or others) to get a second set of eyes <sup>50</sup>. Set a rule: "No merge without at least one thorough review comment or approval." As a manager, follow up on code review discipline – spot-check that reviewers are catching issues (you can even compare notes with Collider's findings to see if reviewers missed something obvious).
- **Mandate Contracts/Test Before Code:** Encourage (or require) that for each feature, the developer (human or AI) produces an API design or test plan first. As a non-coder lead, you can still participate here: review the proposed function list or user acceptance criteria before implementation begins. This not only catches misunderstandings early but sends a signal that *thinking before coding* is expected even with AI. In sprint planning, ask "what will the tests look like?" to reinforce this.

- **Scheduled Refactoring/Quality Sprints:** Put refactoring and cleanup as first-class items in your project schedule. For example, every 4th sprint could include a task like “Consolidate duplicate code” or “Review architecture for drift,” potentially using Collider’s output as a guide. This prevents compounding of AI-generated debt. Track these tasks like any feature – assign owners and celebrate completion.
- **Use Visualization for Communication:** Leverage Collider’s visual topology maps in team meetings. As a manager, you can’t read all the code, but you can absorb a picture of the system structure. If you see a big red node or a weird spiderweb cluster, question it. Ask the team, “*Our architecture diagram didn’t have these many cross-connections – what happened here?*” This encourages the team to use tools and keeps architecture in focus. It also helps you identify when AI might be introducing structural complexity that wasn’t intended.
- **Implement a “Two-Strikes” Rule for AI mistakes:** If an AI introduces a serious issue twice, pause and review the process. For instance, if an AI-generated commit causes a production bug, require a retrospective: how did it slip through? Were tests lacking? Use that as a learning to update your guardrails. This creates accountability – the team knows that while AI can be used, it’s not an excuse for poor quality.
- **Continuous Training and Calibration:** Ensure the AI models or prompts are kept up-to-date with your evolving code standards. From a management perspective, this means allocating time for engineers to fine-tune prompts or model settings when the project’s style changes. If you adopt a new framework or pattern, make sure the AI’s context (like the orientation files) are updated to reflect that. Essentially, treat the AI like a team member who needs onboarding and training when things change.
- **Promote a Culture of Skepticism (Healthy):** Encourage developers to treat AI output as suggestions, not truths. Simple practice: if an AI writes a critical piece of code, have the author explain it in the next stand-up (as if they wrote it themselves). This ensures they fully understand it. Non-coding leads can still foster this by asking questions like “How do we know this code is safe?” or “Did we consider alternatives or edge cases here?” If the answer is “the AI said so,” then push back.

By following this playbook, project leads can create an environment where AI is a productivity booster **without** letting its natural liabilities undermine the project. The combination of automated checks (the “hard guardrails”) and team culture (the “soft guardrails”) will keep development on track. Remember, the goal is to harness AI’s speed **and** maintain human oversight – as one analysis put it, *AI can get you 70% of the way quickly, but the final 30% – the polish, the hardening – demands our intervention* <sup>51</sup>. As a manager, your role is to ensure that final 30% is never skipped.

#### **Sources:**

1. Tenable Research – *FAQ on Vibe Coding* (AI coding risks) <sup>52</sup> 1
2. Nolan Lawson – *How I use AI agents to write code* (experience report on AI mistakes) <sup>26</sup> 7
3. Endor Labs – *Security Vulnerabilities in AI-Generated Code* (hallucinated dependencies, drift) <sup>10</sup> 17
4. Kodus Blog – *Biggest Dangers of AI-Generated Code* (bug patterns, context issues) <sup>12</sup> 2
5. AlterSquare – *Why AI-Generated Code Costs More to Maintain* (duplication, technical debt) <sup>5</sup> 23

6. AlterSquare – *Balancing AI Code with Human Oversight* (recommendations for standards and reviews)

53 25

7. H. Abbasi – *Promise and Pitfalls of AI-Assisted Generation* (skill erosion, oversight) 34 51

8. Collider Documentation – *Standard Model & Antimatter Laws* (8D system, layer rules) 18 39.

---

1 3 6 52 - Blog |

<https://www.tenable.com/blog/frequently-asked-questions-about-vibe-coding>

2 4 12 14 22 27 50 The Biggest Dangers of AI-Generated Code

<https://kodus.io/en/the-biggest-dangers-of-ai-generated-code/>

5 19 23 24 25 30 31 32 33 35 36 37 44 53 Why AI-Generated Code Costs More to Maintain Than Human-Written Code

<https://www.altersquare.io/ai-generated-code-maintenance-costs/>

7 9 21 26 45 How I use AI agents to write code | Read the Tea Leaves

<https://nolanlawson.com/2025/12/22/how-i-use-ai-agents-to-write-code/>

8 THEORY.md

<file://file-7aTzMHSPLhbJJa8Kx9ZyWD>

10 11 13 15 16 17 20 41 The Most Common Security Vulnerabilities in AI-Generated Code | Blog |

Endor Labs

<https://www.endorlabs.com/learn/the-most-common-security-vulnerabilities-in-ai-generated-code>

18 28 29 38 39 40 48 49 SYNTHESIS\_GAP\_IMPLEMENTATION.md

<file://file-7kCfFwGdpnnhg49W5J1z>

34 51 The Promise and Pitfalls of AI-Assisted Code Generation: Balancing Productivity, Control, and Risk |

by Hammad Abbasi | Medium

<https://hammadulhaq.medium.com/the-promise-and-pitfalls-of-ai-assisted-code-generation-balancing-productivity-control-and-risk-80869176202f>

42 43 46 47 TOOL.md

<file://file-UvHTYvmckoPAdqKRgQ5MDB>